



AI-Native Homelab

Private, cost-efficient MLOps you control

What it is • Why it matters • What you get

Powered by ProxmoxMCP + Model Context Protocol

Simon Renauld

Data Engineering Leader • simondatalab.de



The Challenge

- Slow idea → experiment cycles (days, not hours)
- Unpredictable cloud costs for iteration
- Data residency/privacy constraints
- Fragmented tooling across teams

Objective

Make experiments fast, affordable, and compliant — on your own hardware



The Solution

What it is

An AI-managed Proxmox platform (ProxmoxMCP + MCP)

How it works

Agents provision test VMs per PR, run checks, and clean up

For ML

On-demand GPU VMs; training and inference spin up only when needed

Operations

Health checks, logs, and alerts wired into Grafana/MLflow



Example Outcomes

< 2 min

VM provisioning time

50%

fewer manual interventions

20%

lower power/capacity overhead

~5 min

mean incident response



Use Cases

Private RAG

Answer docs/Q&A without sending data to cloud

Analytics to Insight

ETL → model → Grafana dashboard in a day

ML Experimentation

Run 10-20 safe experiments/day on fixed budget

On-prem Inference

Latency-sensitive apps, GPU only when needed

Governed Prototyping

Keep PII/data in-house while iterating fast



Build an AI-Native Platform that Pays Off

- 30-90 day roadmap
- Fixed, transparent budget
- On-prem or hybrid cloud

Simon Renauld

linkedin.com/in/simonrenauld • simondatalab.de